

Dondapati Rahul

Data Engineer

Aurora, IL | (331) 239-0321| RahulD091@outlook.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Data Engineer with 4 years of experience designing, building, and supporting scalable batch and streaming data pipelines across GCP and Azure. Strong expertise in Python, SQL, PySpark, Apache Spark, Snowflake, Databricks, Delta Lake, Airflow, dbt, Kafka, and cloud-based data platforms. Experienced in developing ETL/ELT pipelines, data ingestion frameworks, data lakehouse solutions, streaming workflows, data quality controls, and production-ready data services. Skilled in data modeling, pipeline optimization, orchestration, CI/CD, and cloud infrastructure, with experience supporting financial datasets across investments, portfolios, payments, lending, transactions, and enterprise reporting.

TECHNICAL SKILLS

Programming & Query Languages: Python, SQL, T-SQL, PySpark, Bash, Shell Scripting

Cloud Platforms: Google Cloud Platform (GCP), Microsoft Azure, BigQuery, Google Cloud Storage, Dataproc, Dataflow, Pub/Sub, Cloud Composer, Azure Data Factory, Azure Data Lake Storage Gen2 (ADLS Gen2), Synapse Analytics, Azure Monitor

Data Engineering & Processing: Apache Spark, PySpark, Databricks, Delta Lake, Snowflake, Apache Airflow, dbt, Informatica PowerCenter, ETL/ELT, Batch Processing, Incremental Processing, Data Ingestion, Data Transformation, Metadata-Driven Pipelines, Data Pipeline Orchestration

Streaming & Big Data: Apache Kafka, Apache Flink, Spark Structured Streaming, Apache Beam, Google Dataflow, Pub/Sub, Hadoop, HDFS, Hive, Sqoop

Data Modeling & Architecture: Dimensional Modeling, Star Schema, Snowflake Schema, Fact and Dimension Tables, Slowly Changing Dimensions (SCD), Source-to-Target Mapping, Lakehouse Architecture, Medallion Architecture, Data Warehouse Design

Data Quality & Governance: Great Expectations, dbt Tests, SQL Validation, Data Reconciliation, Source-to-Target Validation, Data Quality Monitoring, Quarantine Tables, Duplicate Detection, Exception Handling, Control-Total Validation, Data Lineage

Databases & Storage: Snowflake, SQL Server, Oracle, BigQuery, Delta Lake, PostgreSQL, Redis, Google Cloud Storage, ADLS Gen2

DevOps & Production Engineering: Git, GitHub Actions, Azure DevOps, CI/CD, Docker, Terraform, Kubernetes, Production Support, Pipeline Monitoring, Dependency Analysis

Security & Data Services: FastAPI, REST APIs, IAM, RBAC, Tokenization, Encryption, Dynamic Data Masking, Secure Data Access

AI & Machine Learning Data Platforms: Vertex AI, MLflow, Embeddings, Vector Search, AI-Ready Data Pipelines, RAG Data Pipelines

PROFESSIONAL EXPERIENCE

Kestra Financial | Data Engineer | USA

Jan 2025 - Present

- Engineered GCP-based data pipelines using Python, PySpark, Dataproc, Cloud Composer, Snowflake, and dbt to integrate investment, portfolio, pricing, asset servicing, and enterprise financial data.
- Developed batch and streaming data workflows using Airflow, Kafka, Pub/Sub, Dataflow, and Flink to process portfolio, transaction, market, and operational data across enterprise platforms.
- Optimized investment and operational reporting pipelines, improving batch completion time by 35% through incremental processing, partition optimization, Snowflake performance tuning, and PySpark optimization.
- Established automated data quality frameworks using Great Expectations, dbt tests, SQL validations, quarantine tables, and source-to-target controls, reducing reconciliation effort by 40%.
- Prepared governed and curated datasets for analytics and AI use cases using Databricks, Delta Lake, Vertex AI embeddings, and MLflow to support enterprise intelligence and data-driven decision-making.
- Built secure data services using FastAPI, Redis, IAM, tokenization, Terraform, and Docker to provide controlled, reliable, and scalable access to enterprise data products.
- Automated testing and deployment across data workflows and production environments using GitHub Actions and CI/CD practices, improving release reliability and operational efficiency.

Virtusa | Data Engineer | INDIA

Aug 2020 - Aug 2023

- Integrated payment, settlement, refund, and chargeback data using SQL and Informatica PowerCenter to deliver standardized datasets for downstream reporting and operational analysis.
- Developed reusable ETL workflows across Oracle, SQL Server, Hadoop, Hive, and Sqoop to support scheduled ingestion, historical data processing, and enterprise data integration.
- Automated file validation, duplicate detection, exception handling, and control-total verification using Python and SQL, reducing payment reconciliation effort by 30%.
- Implemented Azure Data Factory pipelines to ingest lending, credit, repayment, and collection data into Azure Data Lake Storage Gen2 for centralized storage and processing.
- Transformed and standardized customer datasets using PySpark, Azure Databricks, Delta Lake, and Synapse Analytics while maintaining historical records and trusted analytical data layers.
- Designed metadata-driven ETL workflows using configurable source mappings, reusable transformation patterns, and Delta Lake MERGE operations to simplify onboarding of new data sources.
- Investigated production pipeline failures using Azure Monitor, dependency analysis, and source-to-target validation while managing code releases through Git and Azure DevOps.
- Partnered with business and technical teams to resolve data issues, validate pipeline changes, perform user acceptance checks, and support successful production releases.

EDUCATION

Master of Science in Computer Science

Sep 2023 - Jun 2025

City University of Seattle, Seattle, WA, USA

CERTIFICATIONS

- Microsoft Learn – Power BI Data Analyst Learning Path
- Databricks Fundamentals Accreditation
- Snowflake Hands-On Essentials Badge