

Rahul Dondapati

Data Engineer

(425) 520 - 4677 | dondapatirahul19@gmail.com | LinkedIn

Professional Summary

Experienced Data Engineer with 4 years of experience. Specializes in designing, building, and supporting batch and streaming data pipelines across GCP and Azure environments. Skilled in Python, SQL, PySpark, Snowflake, Databricks, Delta Lake, Airflow, Azure Data Factory, dbt, and Kafka. Experienced in data modeling, ETL development, cloud migration, reconciliation, data quality, governance, performance tuning, and production support. Delivers reliable financial data platforms for payments, lending, investment reporting, portfolio analytics, regulatory reporting, and secure business operations across development and production environments.

Technical Skills

Programming & Scripting: Python, Advanced SQL, PySpark, Bash, Shell Scripting, Linux, pandas, JSON, XML, YAML

Cloud Data Platforms: Google Cloud Platform, Microsoft Azure, Google Cloud Storage, BigQuery, Dataproc, Dataflow, Pub/Sub, Cloud Composer, Azure Data Factory, Azure Data Lake Storage Gen2, Azure Synapse Analytics, Azure Event Hubs

Data Warehouses & Lakehouse: Snowflake, Databricks, Delta Lake, BigQuery, Azure Synapse Analytics, Hadoop HDFS, Apache Hive, Medallion Architecture, Data Lakehouse Architecture

ETL, ELT & Orchestration: Apache Airflow, Google Cloud Composer, dbt, Informatica PowerCenter, Azure Data Factory, Apache Sqoop, REST API Integration, Batch Processing, Incremental Loading, Metadata-Driven Pipelines

Big Data & Streaming: Apache Spark, PySpark, Apache Kafka, Apache Flink, Apache Beam, Google Dataflow, Spark Structured Streaming, Pub/Sub, Azure Event Hubs, Event-Time Processing

Databases & Data Modeling: Oracle, SQL Server, Dimensional Modeling, Star Schema, Snowflake Schema, Data Vault, Fact and Dimension Tables, Slowly Changing Dimensions, Source-to-Target Mapping, Security Master Data

Data Quality, Governance & Security: Great Expectations, dbt Tests, Data Reconciliation, Data Validation, Data Contracts, Data Lineage, Duplicate Detection, IAM, RBAC, Encryption, Tokenization, Dynamic Data Masking, Azure Key Vault, Google Secret Manager

DevOps, Monitoring & AI-Ready Data: Git, GitHub, GitHub Actions, Azure DevOps, Docker, Kubernetes, Terraform, CI/CD, Azure Monitor, OpenTelemetry, Vertex AI, MLflow, Feast Feature Store, Embeddings, Vector Search, RAG Data Pipelines, PII Redaction

Work Experience

Charles Schwab | Data Engineer | USA **Jan 2025 - Present**

- Engineered GCP-based lakehouse pipelines using Python, PySpark, Dataproc, Cloud Composer, Snowflake, and dbt to unify brokerage, holdings, pricing, and regulatory datasets.
- Orchestrated batch and streaming workflows with Airflow, Kafka, Pub/Sub, Dataflow, and Flink, ensuring reliable ingestion across portfolio, trade, and market data sources.
- Improved investment-reporting batch completion by 35% by applying incremental processing, partitioning, Snowflake warehouse tuning, and optimized PySpark transformations across regulated data pipelines.
- Reduced portfolio reconciliation effort by 40% by implementing Great Expectations, dbt tests, SQL validations, quarantine tables, and automated source-to-target controls.
- Prepared AI-ready research datasets using Python, Databricks, Delta Lake, Vertex AI embeddings, MLflow, and governed metadata for portfolio insight services.
- Secured and deployed FastAPI data services using Redis, IAM, tokenization, Terraform, Docker, and GitHub Actions, supporting governed research delivery and production reliability.

Persistent Systems | Data Engineer | INDIA **Aug 2020 – Aug 2023**

- Analyzed payment, settlement, refund, and chargeback sources using SQL and Informatica PowerCenter, translating inconsistent records into reliable warehouse-ready financial datasets.
- Constructed reusable ETL workflows across Oracle, SQL Server, Hadoop, Hive, and Sqoop, supporting scheduled ingestion, historical tracking, and daily reconciliation operations.
- Automated manual payment reconciliation effort by 30% by automating file validation, duplicate detection, exception handling, and control-total checks with Python and SQL.
- Designed Azure Data Factory pipelines that ingested lending, credit, repayment, and collection data into Data Lake Storage for controlled downstream processing.
- Integrated PySpark, Azure Databricks, Delta Lake, and Synapse Analytics to standardize customer identifiers, maintain history, and publish trusted lending datasets.
- Accelerated end-to-end pipeline performance by 35% through incremental loading, Delta merge logic, optimized partitions, and metadata-driven orchestration within Azure Data Factory.
- Resolved production failures by reviewing Azure Monitor alerts, validating source-to-target counts, tracing dependencies, and coordinating tested releases through Git and Azure DevOps.

Education

Master of Science in Computer Science | City University of Seattle, Seattle, WA, USA **Sep 2023 - Jun 2025**

Certifications

- Google Data Analytics Professional Certificate – Coursera
- SQL (Intermediate) Certificate – HackerRank